Homework 3

Aline Bessa and Rao Li

March 31, 2019

Question 1

Before computing principal components, let's represent the points as a matrix, with each attribute as a column and each row as a point:

$$X = \begin{bmatrix} -2 & -2 \\ 0 & 0 \\ 2 & 2 \\ -0.5 & 0.5 \\ 0.5 & -0.5 \end{bmatrix}$$

The mean of both columns is zero, so the data is already mean-centered. Now let's compute the covariance matrix associated to X using the standard formula (denominator $= N - 1$):

$$Cov = \begin{bmatrix} \frac{-2^2+2^2+(-0.5)^2+0.5^2}{4} & \frac{-2*-2+2*2+(-0.5*0.5)+0.5*(-0.5)}{4} \\ \frac{-2*-2+2*2+(-0.5*0.5)+0.5*(-0.5)}{4} & \frac{-2^2+2^2+0.5^2+(-0.5)^2}{4} \end{bmatrix}$$

$$Cov = \begin{bmatrix} 2.125 & 1.875 \\ 1.875 & 2.125 \end{bmatrix}$$

Now, let's compute the eigenvalues and eigenvectors associated to the covariance matrix:

$$Cov - \lambda I = \begin{bmatrix} 2.125 - \lambda & 1.875 \\ 1.875 & 2.125 - \lambda \end{bmatrix}$$

$$det(Cov - \lambda I) = (2.125 - \lambda)^2 - 1.875^2$$
$$det(Cov - \lambda I) = 4.515625 - 4.25\lambda + \lambda^2 - 3.515625$$
$$det(Cov - \lambda I) = \lambda^2 - 4.25\lambda + 1$$

The eigenvalues are the solutions to $det(Cov - \lambda I) = 0$:

$$\lambda^2 - 4.25\lambda + 1 = 0$$
$$\lambda = \frac{4.25 \pm \sqrt{(-4.25)^2 - 4 * 1 * 1)}}{2 * 1}$$
$$\lambda = 4 \text{ or } \lambda = 0.25$$

To get the first eigenvactor, using $\lambda = 4$, we do

$$\begin{bmatrix} 2.125 - 4 & 1.875 \\ 1.875 & 2.125 - 4 \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$
$$\begin{bmatrix} -1.875 & 1.875 \\ 1.875 & -1.875 \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$
$$\begin{bmatrix} -1.875x_1 + 1.875x_2 \\ 1.875x_1 - 1.875x_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

That is, $x_1 = x_2$ and they can be any value. So our first eigenvector, which is our first principal component because it is associated to the largest eigenvalue, is

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

As for $\lambda = 0.25$, we have our second principal component:

$$\begin{bmatrix} 2.125 - 0.25 & 1.875 \\ 1.875 & 2.125 - 0.25 \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$
$$\begin{bmatrix} 1.875 & 1.875 \\ 1.875 & 1.875 \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$
$$\begin{bmatrix} 1.875x_1 + 1.875x_2 \\ 1.875x_1 + 1.875x_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

That is, $x_1 = -x_2$ and they can be any value. So our second principal component is:

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

Question 2

To project the points onto the two principle components, we start by creating a matrix with the eigenvectors:

$$E = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

The first and second components correspond, respectively, to the first and second lines of E. We now obtain the projections by multiplying E by X^T:

$$P = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \times \begin{bmatrix} -2 & 0 & 2 & -0.5 & 0.5 \\ -2 & 0 & 2 & 0.5 & -0.5 \end{bmatrix} = \begin{bmatrix} -4 & 0 & 4 & 0 & 0 \\ 0 & 0 & 0 & -1 & 1 \end{bmatrix} =$$

So the projected points are, in the order given in Question 1: $\{(-4, 0), (0, 0), (4, 0), (0, -1), (0, 1)\}$.
CHECK CALCULATIONS!

Question 3

Let's start with the entropy criterion. Suppose $class_1 = A$ and $class_0 = B$. There are then 4 A and 2 B training examples. If we choose feature x_1, we have two subsets of examples: one to which $x_1 = 1$ (S_l) and one to which $x_1 = 0$ (S_r). The entropy for these subsets is

$$
\begin{aligned}
H(S_l) &= -(\frac{3}{3}\log_2\frac{3}{3} + \frac{0}{3}\log_2\frac{0}{3}) \\
H(S_l) &= -1\log_2 1 = 0 \\
H(S_r) &= -(\frac{2}{3}\log_2\frac{2}{3} + \frac{1}{3}\log_2\frac{1}{3}) \\
H(S_r) &= 0.918
\end{aligned}
$$

Finally,

$$H(after) = \frac{|S_l|H(S_l) + |S_r|H(S_r)}{|S_l| + |S_r|} = \frac{3*0 + 3*0.918}{3+3} = 0.459$$

Analogously, for x_2 we have

$$
\begin{aligned}
H(S_l) &= -(\frac{1}{2}\log_2\frac{1}{2} + \frac{1}{2}\log_2\frac{1}{2}) \\
H(S_l) &= 1 \\
H(S_r) &= -(\frac{2}{2}\log_2\frac{2}{2} + \frac{0}{2}\log_2\frac{0}{2}) \\
H(S_r) &= -1\log_2 1 = 0 \\
H(after) &= \frac{|S_l|H(S_l) + |S_r|H(S_r)}{|S_l| + |S_r|} = \frac{4*1 + 2*0}{4+2} = 0.667
\end{aligned}
$$

Analogously, for x_3 we have

$$
\begin{aligned}
H(S_l) &= -(\frac{3}{4}\log_2\frac{3}{4} + \frac{1}{4}\log_2\frac{1}{4}) \\
H(S_l) &= 0.811 \\
H(S_r) &= -(\frac{1}{2}\log_2\frac{1}{2} + \frac{1}{2}\log_2\frac{1}{2}) \\
H(S_r) &= 1 \\
H(after) &= \frac{|S_l|H(S_l) + |S_r|H(S_r)}{|S_l| + |S_r|} = \frac{4*0.811 + 2*1}{4+2} = 0.874
\end{aligned}
$$

Finally,for x_4 we have

$$
\begin{aligned}
&H(S_l) = -(\frac{1}{2}\log_2\frac{1}{2} + \frac{1}{2}\log_2\frac{1}{2})\\
&H(S_l) = 1\\
&H(S_r) = -(\frac{4}{4}\log_2\frac{4}{4} + \frac{0}{4}\log_2\frac{0}{4})\\
&H(S_r) = -1\log_2 1 = 0\\
&H(after) = \frac{|S_l|H(S_l) + |S_r|H(S_r)}{|S_l| + |S_r|} = \frac{4*1+2*0}{4+2} = 0.667
\end{aligned}
$$

Because we want to minimize $H(after)$ to find the best split, x_1 will be chosen for the root.

Now let's use the Gini criterion, using S_r and S_l as defined above for the different x features. For x_1, we have

$$
\begin{aligned}
&G(S_l) = 1 - 1^2 = 0\\
&G(S_r) = 1 - \left(\frac{1}{3}\right)^2 - \left(\frac{2}{3}\right)^2 = \frac{4}{9}\\
&G(S) = \frac{1}{2}*0 + \frac{1}{2}*\frac{4}{9} = 0.222
\end{aligned}
$$

For x_2, we have

$$
\begin{aligned}
&G(S_l) = 1 - \left(\frac{1}{2}\right)^2 - \left(\frac{1}{2}\right)^2 = \frac{1}{2}\\
&G(S_r) = 1 - 1^2 = 0\\
&G(S) = \frac{2}{3}*\frac{1}{2} + \frac{1}{3}*0 = 0.333
\end{aligned}
$$

For x_3, we have

$$
\begin{aligned}
&G(S_l) = 1 - \left(\frac{3}{4}\right)^2 - \left(\frac{1}{4}\right)^2 = \frac{3}{8}\\
&G(S_r) = 1 - \left(\frac{1}{2}\right)^2 - \left(\frac{1}{2}\right)^2 = \frac{1}{2}\\
&G(S) = \frac{2}{3}*\frac{3}{8} + \frac{1}{3}*\frac{1}{2} = 0.417
\end{aligned}
$$

Finally, for x_4 we have

$$
\begin{aligned}
&G(S_l) = 1 - \left(\frac{1}{2}\right)^2 - \left(\frac{1}{2}\right)^2 = \frac{1}{2}\\
&G(S_r) = 1 - 1^2 = 0\\
&G(S) = \frac{2}{3}*\frac{1}{2} + \frac{1}{3}*0 = 0.333
\end{aligned}
$$

Because the Gini criterion calculates how frequently a randomly chosen element will be wrongly identified, we want to minimize it to find the best split. Consequently, x_1 will be chosen for the root.

Now let's use the Misclassification criterion, using S_r and S_l as defined above for the different x features. For x_1, we have

$$
\begin{aligned}
J(S_l) &= 0 \\
J(S_r) &= 1 \\
J(S) &= 0 + 1 = 1
\end{aligned}
$$

For x_2, we have

$$
\begin{aligned}
J(S_l) &= 2 \\
J(S_r) &= 0 \\
J(S) &= 2 + 0 = 2
\end{aligned}
$$

For x_3, we have

$$
\begin{aligned}
J(S_l) &= 1 \\
J(S_r) &= 1 \\
J(S) &= 1 + 1 = 2
\end{aligned}
$$

Finally, for x_4 we have

$$
\begin{aligned}
J(S_l) &= 2 \\
J(S_r) &= 0 \\
J(S) &= 2 + 0 = 2
\end{aligned}
$$

Because this criterion should minimize the number of points that are incorrectly classified, x_1 will be chosen for the root.

Question 4

Let the discriminant functions be

$$
\begin{aligned}
g_1(x_1, x_2) &= 5x_2 + 3x_1 - 4 \\
g_2(x_1, x_2) &= -3x_2 + 2x_1 - 6
\end{aligned}
$$

We assign an example (x_1, x_2) to class C_1 when $g_1(x_1, x_2) > g_2(x_1, x_2)$, that is

$$
\begin{aligned}
&g_1(x_1, x_2) > g_2(x_1, x_2) \\
&5x_2 + 3x_1 - 4 > -3x_2 + 2x_1 - 6 \\
&8x_2 - x_1 + 2 > 0 \\
&g(x_1, x_2) = 8x_2 - x_1 + 2
\end{aligned}
$$

So if $g(x_1, x_2) > 0$, the example is assigned to class C_1; otherwise, to class C_2.

Question 5

(a) When there are two classes, the maximum entropy occurs when they are equaly likely, i.e., when *half* of the examples are positive and *half* are negative. The closer the proportions are to $\frac{1}{2}$, the higher the entropy. In the first dataset, we have that the proportions for positive and negative class are, respectively, $\frac{4}{9}$ and $\frac{5}{9}$. Consequently, the difference between these proportions and $\frac{1}{2}$ are the same, and can be calculated as

$$|\frac{4}{9} - \frac{1}{2}| = \frac{|2 * 4 - 9 * 1|}{18} = \frac{1}{18}$$

As for the second dataset, the proportions for positive and negative class are, respectively, $\frac{1}{3}$ and $\frac{2}{3}$. The difference between these proportions and $\frac{1}{2}$ are the same, and can be calculates as

$$|\frac{1}{3} - \frac{1}{2}| = \frac{|2 * 1 - 3 * 1|}{6} = \frac{1}{6}$$

Given that the difference for the first dataset is smaller, its entropy is higher (this dataset has more *impurity*). In other words, the entropy for the dataset with 4 positive and 5 negative examples is higher.

(b) First, let's compute the entropy of the entire dataset, namely S:

$$\begin{aligned}
&Entropy(S) = -(\frac{3}{7}\log_2\frac{3}{7} + \frac{3}{7}\log_2\frac{3}{7})\\
&Entropy(S) = 0.98522813603425152
\end{aligned}$$

Now, let's compute the entropy associated to the examples where $x_1 = F$ and where $x_1 = F$.

$$\begin{aligned}
&Entropy(S_{x_1=F}) = -(\frac{2}{4}\log_2\frac{2}{4} + \frac{2}{4}\log_2\frac{2}{4})\\
&Entropy(S_{x_1=F}) = 1.0\\
&Entropy(S_{x_1=T}) = -(\frac{1}{3}\log_2\frac{1}{3} + \frac{2}{3}\log_2\frac{2}{3})\\
&Entropy(S_{x_1=T}) = 0.91829583405448956
\end{aligned}$$

Consequently, the second term of the Information Gain formula is

$$\begin{aligned}
&Y = \sum_{v\in\{F,T\}} \frac{|S_{x_1=v}|}{|S|} Entropy(S_{x_1=v})\\
&Y = \frac{4}{7}1.0 + \frac{3}{7}0.91829583405448956\\
&Y = 0.9649839288804954
\end{aligned}$$

The final value for x_1 is thus

$$\begin{aligned}
&Information - Gain(S) = 0.98522813603425152 - 0.9649839288804954\\
&Information - Gain(S) = 0.020244207153756077
\end{aligned}$$

(c) WILL TYPE SOON

(d) First, let's compute $H(Y)$.

$$H(Y) = -(P[Y=+]\log_2 P[Y=+] + P[Y=-]\log_2 P[Y=-])$$
$$H(Y) = -(\frac{3}{7}\log_2\frac{3}{7} + \frac{4}{7}\log_2\frac{4}{7})$$
$$H(Y) = 0.98522813603425152$$

Now, let's compute $H(Y|X)$.

$$H(Y|X) = \sum_x P[X=x] * \Big(\sum_y -P[Y=y|X=x] * \log_2 P[Y=y|X=x]\Big)$$
$$H(Y|X) = \frac{4}{7}(-\frac{1}{2}\log_2\frac{1}{2} - \frac{1}{2}\log_2\frac{1}{2}) + \frac{3}{7}(-\frac{1}{3}\log_2\frac{1}{3} - \frac{2}{3}\log_2\frac{2}{3})$$
$$H(Y|X) = \frac{4}{7} * 1.0 + \frac{3}{7} * 0.91829583405448956$$
$$H(Y|X) = 0.9649839288804954$$

Finally,

$$H(Y) - H(Y|X) = 0.98522813603425152 - 0.9649839288804954$$
$$H(Y) - H(Y|X) = 0.020244207153756077$$

(e) Using the entropy formula for a dataset S, we have that

$$Entropy(S) = -\sum_{i \in z}\frac{N_i}{N}\log_2\frac{N_i}{N}$$

If each label is equally likely, we can write $\frac{N_i}{N} = \frac{1}{|z|}$ for any i. Consequently,

$$Entropy(S) = -\sum_{i \in z}\frac{1}{|z|}\log_2\frac{1}{|z|}$$
$$Entropy(S) = -|z|\frac{1}{|z|}\log_2\frac{1}{|z|}$$
$$Entropy(S) = -\log_2\frac{1}{|z|}$$
$$Entropy(S) = \log_2|z|$$

where $|z|$ is the number of different labels.